

HW4

November 29, 2021

1 CSE 252A Computer Vision I Fall 2021 - Assignment 4

1.1 Instructor: Ben Ochoa

- Assignment Published On: **Wed, November 17, 2021.**
- Due On: **Wed, December 1, 2021 11:59 PM (Pacific Time).**

1.2 Instructions

- This assignment must be completed **individually**. For more details, please follow the Academic Integrity Policy and Collaboration Policy on [Canvas](#).
- All solutions must be written in this notebook.
 - If it includes the theoretical problems, you **must** write your answers in Markdown cells (using LaTeX when appropriate).
 - Programming aspects of the assignment must be completed using Python in this notebook.
- You may use Python packages (such as NumPy and SciPy) for basic linear algebra, but you may not use packages that directly solve the problem.
 - If you are unsure about using a specific package or function, then ask the instructor and/or teaching assistants for clarification.
- You must submit this notebook exported as a PDF that contains separate pages. You must also submit this notebook as .ipynb file.
 - Submit both files (.pdf and .ipynb) on Gradescope.
 - **You must mark the PDF pages associated with each question in Gradescope. If you fail to do so, we may dock points.**
- It is highly recommended that you begin working on this assignment early.
- **Late Policy:** Assignments submitted late will receive a 15% grade reduction for each 12 hours late (i.e., 30% per day). Assignments will not be accepted 72 hours after the due date. If you require an extension (for personal reasons only) to a due date, you must request one as far in advance as possible. Extensions requested close to or after the due date will only be granted for clear emergencies or clearly unforeseeable circumstances.

1.3 Problem 1: Machine Learning [28 pts]

In this problem, you will implement several machine learning solutions for computer vision problems.

1.3.1 Problem 1.1: Initial Setup

We will use [Scikit-learn \(Sklearn\)](#) module in for this problem. It is the most useful and robust library for machine learning in Python. It provides a selection of efficient tools for machine learning and statistical modeling including classification, regression, clustering and dimensionality reduction via a consistence interface in Python. This library, which is largely written in Python, is built upon NumPy, SciPy and Matplotlib.

Get started by installing the Sklearn module.

```
[204]: import sklearn
sklearn.__version__
```

```
[204]: '0.24.1'
```

1.3.2 Problem 1.2: Download MNIST data [3 pts]

The [MNIST database](#) (Modified National Institute of Standards and Technology database) is a well-known dataset consisting of 28x28 grayscale images of handwritten digits. For this problem, we will use Sklearn to do machine learning classification on the MNIST database.

Sklearn provides a subset of MNIST database with 8x8 pixel images of digits. The images attribute of the dataset stores 8x8 arrays of grayscale values for each image. The target attribute of the dataset stores the digit each image represents. Complete `plot_mnist_sample()` to plot a 2x5 figure, each grid lies a sample image from a category. The following image gives an example:

```
[205]: import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
```

```
[206]: # Download MNIST Dataset from Sklearn
digits = datasets.load_digits()

# Print to show there are 1797 images (8 by 8 images for a dimensionality of 64)
print("Image Data Shape" , digits.data.shape)

# Print to show there are 1797 labels (integers from 0-9)
print("Label Data Shape", digits.target.shape)
```

```
Image Data Shape (1797, 64)
```

```
Label Data Shape (1797,)
```

```
[207]: import random
def plot_mnist_sample():
    """
    This function plots a sample image for each category,
    The result is a figure with 2x5 grid of images.

    """
    plt.figure(figsize=(14,7))
```

```

for i in range(10):
    plotted = False
    while not plotted:
        n = random.randint(0,digits.target.shape[0])
        if digits.target[n] == i:
            plt.subplot(2,5,i+1)
            plt.axis('off')
            plt.title(f"Label: {i}")
            plt.imshow(digits.data[n].reshape([8,8]),cmap="gray")
            plotted = True


```

```

[208]: # PLOT CODE: DO NOT CHANGE
# This code is for you to plot the results.

plot_mnist_sample()

```



1.3.3 Problem 1.3: Recognizing hand-written digits with Sklearn [5 pts]

One of the most amazing things about Sklearn library is that it provides an easy pattern for you to call different models. In this part, we will get some experience with several classifiers in Sklearn. You will complete `LogisticRegressionClassifier` and `kNNClassifier`.

```

[209]: # DO NOT CHANGE
#### Some helper functions are given below####
def DataBatch(data, label, batchsize, shuffle=True):
    """
    This function provides a generator for batches of data that

```

```

yields data (batchsize, 3, 32, 32) and labels (batchsize)
if shuffle, it will load batches in a random order
"""

n = data.shape[0]
if shuffle:
    index = np.random.permutation(n)
else:
    index = np.arange(n)
for i in range(int(np.ceil(n/batchsize))):
    inds = index[i*batchsize : min(n,(i+1)*batchsize)]
    yield data[inds], label[inds]

def test(testData, testLabels, classifier):
    """
    Call this function to test the accuracy of a classifier
    """

    batchsize=50
    correct=0.
    for data,label in DataBatch(testData,testLabels,batchsize,shuffle=False):
        prediction = classifier(data)
        correct += np.sum(prediction==label)
    return correct/testData.shape[0]*100

```

```

[210]: # DO NOT CHANGE
# Split data into 50% train and 50% test subsets
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    digits.images.reshape((len(digits.images), -1)), digits.target, test_size=0.
    ↪5, shuffle=False)

```

```

[211]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

class RandomClassifier():
    """
    This is a sample classifier.
    given an input it outputs a random class
    """

    def __init__(self, classes=10):
        self.classes=classes
    def __call__(self, x):
        return np.random.randint(self.classes, size=x.shape[0])

class LogisticRegressionClassifier():
    def __init__(self, sol='liblinear'):
        self.classifier = LogisticRegression(solver=sol)

```

```

def train(self, trainData, trainLabels):
    self.classifier.fit(trainData, trainLabels)

def __call__(self, x):
    return self.classifier.predict(x)

class kNNClassifier():
    def __init__(self, k=3):
        """
        k is the number of neighbors involved in voting
        """
        self.classifier = KNeighborsClassifier(n_neighbors=k)

    def train(self, trainData, trainLabels):
        self.classifier.fit(trainData, trainLabels)

    def __call__(self, x):
        """
        this method should take a batch of images and return a batch of
        → predictions
        """
        return self.classifier.predict(x)

```

```

[212]: # TEST CODE: DO NOT CHANGE
randomClassifierX = RandomClassifier()
print ('Random classifier accuracy: %f'%test(X_test, y_test, randomClassifierX))

```

Random classifier accuracy: 9.788654

```

[213]: # TEST CODE: DO NOT CHANGE
# TEST LogisticRegressionClassifier

lrClassifierX = LogisticRegressionClassifier()
lrClassifierX.train(X_train, y_train)
print ('Logistic Regression Classifier classifier accuracy: %f'%test(X_test,
→ y_test, lrClassifierX))

```

Logistic Regression Classifier classifier accuracy: 91.657397

```

[214]: # TEST CODE: DO NOT CHANGE
# TEST kNNClassifier

knnClassifierX = kNNClassifier()
knnClassifierX.train(X_train, y_train)

```

```
print ('KNN Classifier classifier accuracy: %f'%test(X_test, y_test,
→knnClassifierX))
```

KNN Classifier classifier accuracy: 96.329255

1.3.4 Problem 1.4: Confusion Matrix [5 pts]

A confusion matrix is a table that is often used to describe the performance of a classification model (or “classifier”) on a set of test data for which the true values are known. Here you will implement a function that computes the confusion matrix for a classifier. The matrix (M) should be $n \times n$ where n is the number of classes. Entry $M[i, j]$ should contain the fraction of images of class i that was classified as class j . The following example plots confusion matrix for the RandomClassifier, your task is to plot the results for LogisticRegressionClassifier and kNNClassifier.

```
[227]: from tqdm import tqdm

def Confusion(testData, testLabels, classifier):
    batchsize=50
    correct=0
    M=np.zeros((10,10))
    num=testData.shape[0]/batchsize
    count=0
    acc=0

    for data,label in tqdm(DataBatch(testData,testLabels,batchsize,shuffle=False),total=len(testData)/
→/batchsize):
        pred_label = classifier(data)
        for b in range(len(pred_label)):
            M[label[b]][pred_label[b]] += 1
            if label[b] == pred_label[b]:
                acc += 1
    for i in range(10):
        M[i] = M[i]/np.sum(M[i])

    return M,acc*100.0/len(testData)


def VisualizeConfussion(M):
    plt.figure(figsize=(14, 6))
    plt.imshow(M)
    plt.show()
    print(np.round(M,2))
```

```
[228]: # TEST/PLOT CODE: DO NOT CHANGE
# TEST LogisticRegressionClassifier

M,acc = Confusion(X_test, y_test, lrClassifierX)
```

```
VisualizeConfussion(M)
```

```
18it [00:00, 3949.23it/s]
```




```
[[0.97 0. 0. 0. 0.01 0.01 0.01 0. 0. 0. ]
 [0. 0.89 0. 0.01 0. 0. 0.01 0. 0.02 0.07]
 [0.02 0. 0.98 0. 0. 0. 0. 0. 0. 0. ]
 [0. 0.01 0. 0.82 0. 0.07 0. 0.02 0.05 0.02]
 [0.01 0.02 0. 0. 0.93 0. 0.01 0. 0. 0.02]
 [0. 0.03 0. 0. 0. 0.9 0.02 0. 0. 0.04]
 [0. 0. 0.01 0. 0. 0. 0.99 0. 0. 0. ]
 [0. 0.01 0. 0. 0. 0.06 0. 0.89 0.01 0.03]
 [0. 0.06 0.01 0. 0. 0.03 0.01 0. 0.88 0.01]
 [0.02 0.01 0. 0. 0. 0.02 0. 0. 0.02 0.92]]
```

```
[229]: # TEST/PLOT CODE: DO NOT CHANGE
# TEST kNNclassifier
```

```
M,acc = Confusion(X_test, y_test, knnClassifierX)
VisualizeConfussion(M)
```

18it [00:00, 198.14it/s]



```
[[0.99 0.  0.  0.  0.01 0.  0.  0.  0.  0. ]
 [0.  0.99 0.  0.  0.  0.  0.  0.  0.01 0. ]
 [0.01 0.  0.94 0.05 0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.  0.92 0.  0.01 0.  0.02 0.03 0.01]
 [0.  0.  0.  0.  0.93 0.  0.  0.  0.  0.07]
 [0.  0.  0.  0.  0.  0.97 0.01 0.  0.  0.02]
 [0.  0.  0.  0.  0.  0.  1.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.  1.  0.  0. ]
 [0.  0.02 0.01 0.01 0.  0.  0.  0.  0.95 0. ]
 [0.  0.  0.  0.03 0.  0.02 0.  0.  0.01 0.93]]
```


1.3.5 Problem 1.5: K-Nearest Neighbors (KNN) [7 pts]

For this problem, you will complete a simple kNN classifier without Sklearn. The distance metric is Euclidean distance (L2 norm) in pixel space. k refers to the number of neighbors involved in voting on the class.

```
[230]: from scipy.spatial import distance

class kNNClassifierManual():
    def __init__(self, k=3):
        self.k=k

    def train(self, trainData, trainLabels):
        self.X_train = trainData
        self.y_train = trainLabels

    def __call__(self, X):
        def get_majority_vote(distances, labels, K):
            # Gathering votes for the K-nearest candidates
            sorted_index_by_distance = np.argsort(np.array(distances), axis=0)
            majority_vote = [0]*10
            majority_index_value = [0]*10
            for k in range(K):
                majority_vote[int(labels[sorted_index_by_distance[k]])] += 1
                majority_index_value[int(labels[sorted_index_by_distance[k]])] += k

            # Selecting the candidate with highest number of votes
            cand_inx = majority_vote.index(max(majority_vote))

            return cand_inx

        # Iterating through every test image
        output = []
        for i in range(X.shape[0]):
            distance_vec = []
            for j in range(self.y_train.shape[0]):
                d = distance.euclidean(X[i], self.X_train[j])
                distance_vec.append(d)

            # Using k nearest clusters to get prediction
            prediction = get_majority_vote(distance_vec, self.y_train, self.k)
            output.append(prediction)
        return output

[231]: # TEST/PLOT CODE: DO NOT CHANGE
# TEST kNNClassifierManual

knnClassifierManualX = kNNClassifierManual()
```

```
knnClassifierManualX.train(X_train, y_train)
print ('KNN classifier accuracy: %f'%test(X_test, y_test, knnClassifierManualX))
```

KNN classifier accuracy: 96.329255

1.3.6 Problem 1.6: Principal Component Analysis (PCA) K-Nearest Neighbors (KNN) [8 pts]

Here you will implement a simple KNN classifier in PCA space (for $k=3$ and 25 principal components). You should implement PCA yourself using `svd` (you may not use `sklearn.decomposition.PCA` or any other package that directly implements PCA transformations)

Is the testing time for PCA KNN classifier more or less than that for KNN classifier? Comment on why it differs if it does.

```
[232]: class PCAKNNClassifier():
    def __init__(self, components=25, k=3):
        # components = number of principal components
        # k is the number of neighbors involved in voting
        self.k = k
        self.components = components
        self.classifier = knnClassifierManual()

    def train(self, trainData, trainLabels):
        _, _, V = np.linalg.svd(np.cov(trainData.T))
        self.W = V[:self.components].T
        trainData_proj = np.dot(trainData, self.W)
        self.classifier.train(trainData_proj, trainLabels)

    def __call__(self, x):
        # this method should take a batch of images
        # and return a batch of predictions
        x_pca = np.dot(x, self.W)
        return self.classifier(x_pca)

# test your classifier with only the first 100 training examples (use this
# while debugging)
pcaknnClassifierX = PCAKNNClassifier()
pcaknnClassifierX.train(X_train[:100], y_train[:100])
print ('PCA KNN classifier accuracy: %f'%test(X_test, y_test, pcaknnClassifierX))
```

KNN classifier accuracy: 85.539488


```
[233]: # test your classifier with all the training examples
pcaknnClassifier = PCAKNNClassifier()
pcaknnClassifier.train(X_train, y_train)
# display confusion matrix for your PCA KNN classifier with all the training
→ examples
```

```
M_pca,acc_pca = Confusion(X_test, y_test, pcaknnClassifier)

print ('PCA KNN classifier accuracy: %f'%acc_pca)
VisualizeConfussion(M_pca)
```

18it [00:07, 2.40it/s]

PCA KNN classifier accuracy: 95.773081



```
[[0.99 0. 0. 0. 0.01 0. 0. 0. 0. 0. ]
 [0. 0.97 0.01 0. 0. 0. 0. 0. 0.01 0.01]
 [0.01 0. 0.94 0.05 0. 0. 0. 0. 0. 0. ]
 [0. 0. 0. 0.9 0. 0.02 0. 0.03 0.04 0. ]
 [0. 0. 0. 0. 0.95 0. 0. 0. 0. 0.05]
 [0. 0. 0. 0. 0. 0.98 0.01 0. 0. 0.01]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. ]
 [0. 0.01 0.01 0. 0. 0. 0. 0.98 0. 0. ]
```

```
[0.  0.02 0.  0.02 0.  0.  0.  0.01 0.94 0. ]
[0.  0.  0.  0.04 0.  0.02 0.  0.  0.  0.93]]
```

The PCA runs a bit faster, but I would have expected a larger gap as the PCA reduces the dimensionality of the points.

1.4 Problem 2: Deep learning [28 pts]

1.4.1 Problem 2.1 Initial setup [1 pts]

Follow the directions on <https://pytorch.org/get-started/locally/> to install Pytorch on your computer.

Note: You will not need GPU support for this assignment so don't worry if you don't have one. Furthermore, installing with GPU support is often more difficult to configure so it is suggested that you install the CPU only version. TA's will not provide any support related to GPU or CUDA.

Run the torch import statements below to verify your instalation.

```
[234]: import torch.nn as nn
import torch.nn.functional as F
import torch
from torch.autograd import Variable

x = torch.rand(5, 3)
print(x)
```

```
tensor([[0.1473, 0.2302, 0.6622],
        [0.4620, 0.2670, 0.8890],
        [0.3864, 0.5406, 0.8807],
        [0.6539, 0.9352, 0.2121],
        [0.4277, 0.1129, 0.2121]])
```

In this problem, we will use the full dataset of MNIST database with 28x28 pixel images of digits.

Download the MNIST data from <http://yann.lecun.com/exdb/mnist/>.

Download the 4 zipped files, extract them into one folder, and change the variable 'path' in the code below. (Code taken from <https://gist.github.com/akesling/5358964>)

Plot one random example image corresponding to each label from training data.

```
[691]: import os
import struct

# Change path as required
path = "./mnist/"

def read(dataset = "training", datatype='images'):
    """
    Python function for importing the MNIST data set. It returns an iterator
    of 2-tuples with the first element being the label and the second element
```

being a numpy.uint8 2D array of pixel data for the given image.
"""

```
if dataset == "training":
    fname_img = os.path.join(path, 'train-images-idx3-ubyte')
    fname_lbl = os.path.join(path, 'train-labels-idx1-ubyte')
elif dataset == "testing":
    fname_img = os.path.join(path, 't10k-images-idx3-ubyte')
    fname_lbl = os.path.join(path, 't10k-labels-idx1-ubyte')

# Load everything in some numpy arrays
with open(fname_lbl, 'rb') as flbl:
    magic, num = struct.unpack(">II", flbl.read(8))
    lbl = np.fromfile(flbl, dtype=np.int8)


with open(fname_img, 'rb') as fimg:
    magic, num, rows, cols = struct.unpack(">IIII", fimg.read(16))
    img = np.fromfile(fimg, dtype=np.uint8).reshape(len(lbl), rows, cols)

if(datatype=='images'):
    get_data = lambda idx: img[idx]
elif(datatype=='labels'):
    get_data = lambda idx: lbl[idx]

# Create an iterator which returns each image in turn
for i in range(len(lbl)):
    yield get_data(i)

X_train=np.array(list(read('training','images'))))
y_train=np.array(list(read('training','labels'))))
X_test=np.array(list(read('testing','images'))))
y_test=np.array(list(read('testing','labels'))))

plt.figure(figsize=(14,7))
for i in range(10):
    plotted = False
    while not plotted:
        n = random.randint(0,X_train.shape[0])
        if y_train[n] == i:
            plt.subplot(2,5,i+1)
            plt.axis('off')
            plt.title(f"Label: {i}")
            plt.imshow(X_train[n].reshape([28,28]), cmap="gray")
            plotted = True
```



1.4.2 Problem 2.2: Training with PyTorch [8 pts]

Below is some helper code to train your deep networks. Complete the train function for DNN below. You should write down the training operations in this function. That means, for a batch of data you have to initialize the gradients, forward propagate the data, compute error, do back propagation and finally update the parameters. This function will be used in the following questions with different networks. You can look at https://pytorch.org/tutorials/beginner/pytorch_with_examples.html for reference.

```
[692]: # base class for your deep neural networks. It implements the training loop
        ↪(train_net).

import torch.nn.init
import torch.optim as optim
from torch.autograd import Variable
from torch.nn.parameter import Parameter
from tqdm import tqdm
from scipy.stats import truncnorm

class DNN(torch.nn.Module):
    def __init__(self, in_features=28*28, classes=10):
        super(DNN, self).__init__()
        pass

    def forward(self, x):
        raise NotImplementedError
```

```

def train_net(self, X_train, y_train, epochs=1, batchSize=50):
    # convert train and test data to tensors
    X_train = torch.Tensor(X_train)
    y_train = torch.LongTensor(y_train)
    # initialize loss function and optimizer for use in each batch
    loss_fn = torch.nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(self.params, lr=0.8*1e-3)
    for epoch in range(epochs):
        for data, label in
→tqdm(DataBatch(X_train, y_train, batchSize, shuffle=False), total=len(X_train)//
→batchSize):
            output = self.forward(data)
            loss = loss_fn(output, label)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

    def __call__(self, x):
        inputs = Variable(torch.FloatTensor(x))
        prediction = self.forward(inputs)
        return np.argmax(prediction.data.cpu().numpy(), 1)

# helper function to get weight variable
def weight_variable(shape):
    initial = torch.Tensor(truncnorm.rvs(-1/0.01, 1/0.01, scale=0.01,
→size=shape))
    return Parameter(initial, requires_grad=True)

# helper function to get bias variable
def bias_variable(shape):
    initial = torch.Tensor(np.ones(shape)*0.1)
    return Parameter(initial, requires_grad=True)

```

```

[693]: # example linear classifier - input connected to output
# you can take this as an example to learn how to extend DNN class
class LinearClassifier(DNN):
    def __init__(self, in_features=28*28, classes=10):
        super(LinearClassifier, self).__init__()
        # in_features=28*28
        self.weight1 = weight_variable((classes, in_features))
        self.bias1 = bias_variable((classes))
        self.params = [self.weight1, self.bias1]

    def forward(self, x):
        # linear operation
        x = torch.Tensor(x)

```

```

        y_pred = torch.addmm(self.bias1, x.view(list(x.size())[0], -1), self.
→weight1.t())
        return y_pred

```

```

X_train=np.float32(np.expand_dims(X_train,-1))/255
X_train=X_train.transpose((0,3,1,2))

```

```

X_test=np.float32(np.expand_dims(X_test,-1))/255
X_test=X_test.transpose((0,3,1,2))

```

```

[686]: # test the example linear classifier (note you should get around 90% accuracy
# for 10 epochs and batchsize 50)
linearClassifier = LinearClassifier()
linearClassifier.train_net(X_train, y_train, epochs=10)

print ('Linear classifier accuracy: %f'%test(X_test, y_test, linearClassifier))

```

```

100%|| 1200/1200 [00:00<00:00, 1617.14it/s]
100%|| 1200/1200 [00:00<00:00, 1755.42it/s]
100%|| 1200/1200 [00:01<00:00, 1191.90it/s]
100%|| 1200/1200 [00:00<00:00, 1854.94it/s]
100%|| 1200/1200 [00:00<00:00, 1335.65it/s]
100%|| 1200/1200 [00:00<00:00, 1876.12it/s]
100%|| 1200/1200 [00:00<00:00, 1899.59it/s]
100%|| 1200/1200 [00:00<00:00, 2033.49it/s]
100%|| 1200/1200 [00:00<00:00, 1937.25it/s]
100%|| 1200/1200 [00:00<00:00, 1932.28it/s]

```

Linear classifier accuracy: 92.660000

```

[687]: # display confusion matrix
M_lin,acc_lin = Confusion(X_test, y_test, linearClassifier)
print ('MLP classifier accuracy: %f'%acc_lin)
VisualizeConfussion(M_lin)


```

```

100%|| 200/200 [00:00<00:00, 1366.52it/s]

```

MLP classifier accuracy: 92.660000




```
[[0.98 0.  0.  0.  0.  0.01 0.01 0.  0.  0. ]
 [0.  0.98 0.  0.  0.  0.  0.  0.  0.01 0. ]
 [0.01 0.01 0.89 0.02 0.01 0.  0.01 0.01 0.04 0. ]
 [0.  0.  0.02 0.91 0.  0.03 0.  0.01 0.02 0.01]
 [0.  0.  0.  0.  0.92 0.  0.01 0.  0.01 0.05]
 [0.01 0.  0.  0.03 0.01 0.89 0.01 0.01 0.03 0.01]
 [0.01 0.  0.01 0.  0.01 0.01 0.95 0.  0.  0. ]
 [0.  0.01 0.02 0.  0.  0.  0.  0.92 0.  0.04]
 [0.01 0.01 0.01 0.02 0.01 0.03 0.01 0.01 0.89 0.01]
 [0.01 0.01 0.  0.01 0.02 0.01 0.  0.02 0.  0.93]]
```

1.4.3 Problem 2.3: Single Layer Perceptron [3 pts]

The simple linear classifier implemented in the cell already performs quite well. Plot the filter weights corresponding to each output class (weights, not biases) as images. (Normalize weights to lie between 0 and 1 and use color maps like 'inferno' or 'plasma' for good results). Comment on what the weights look like and why that may be so.

```
[520]: # Plot filter weights corresponding to each class, you may have to reshape them
        ↳ to make sense out of them
        # linearClassifier.weight1.data will give you the first layer weights
def normalize(x):
    return (x-x.min())/(x.max()-x.min())

weights = linearClassifier.weight1.data
plt.figure(figsize=(16,8))
for i in range(weights.shape[0]):
    weight = weights[i,:]
    # normalize weights
    weight = normalize(weight)
    plt.subplot(2,5,i+1)
    plt.axis('off')
    plt.title(f"Label: {i}")
    plt.imshow(weight.reshape((28,28)), cmap="inferno")
```



Comments on weights We observe that the perceptrons look fairly similar to the actual numbers they are labeled as. We learn the perceptrons to activate on the different numbers only as we take the dot product between the input image and a bias.

1.4.4 Problem 2.4: Multi Layer Perceptron (MLP) [8 pts]

Here you will implement an MLP. The MLP should consist of 2 layers (matrix multiplication and bias offset) that map to the following feature dimensions:

- 28x28 -> hidden (100)
- hidden -> classes

- The hidden layer should be followed with a ReLU nonlinearity. The final layer should not have a nonlinearity applied as we desire the raw logits output.
- The final output of the computation graph should be stored in `self.y` as that will be used in the training.

Display the confusion matrix and accuracy after training. Note: You should get ~ 97 % accuracy for 10 epochs and batch size 50.

Plot the filter weights corresponding to the mapping from the inputs to the first 10 hidden layer outputs (out of 100). Do the weights look similar to the weights plotted in the previous problem? Why or why not?

```
[688]: class MLPClassifier(DNN):
    def __init__(self, in_features=28*28, classes=10, hidden=100):
        super(MLPClassifier, self).__init__()
        #Init first layer
        self.weight1 = weight_variable((hidden,in_features))
        self.bias1 = bias_variable((hidden))
        #Non-linearity
        self.ReLU = torch.nn.ReLU()
        #Init second layer
        self.weight2 = weight_variable((classes,hidden))
        self.bias2 = bias_variable((classes))
        self.params = [self.weight1,self.bias1,self.weight2,self.bias2]

    def forward(self, x):
        x = torch.Tensor(x)
        # Input layer
        out = torch.addmm(self.bias1, x.view(list(x.size())[0], -1), self.
→weight1.t())
        # Activation
        out = self.ReLU(out)
        # Output layer
        out = torch.addmm(self.bias2, out.view(list(x.size())[0], -1), self.
→weight2.t())
        return out

mlpClassifier = MLPClassifier()
mlpClassifier.train_net(X_train, y_train, epochs=10, batchSize=5)
```

```
100%|| 12000/12000 [00:13<00:00, 917.23it/s]
100%|| 12000/12000 [00:13<00:00, 887.12it/s]
100%|| 12000/12000 [00:14<00:00, 830.74it/s]
100%|| 12000/12000 [00:14<00:00, 848.24it/s]
100%|| 12000/12000 [00:14<00:00, 813.48it/s]
```

```

100%|| 12000/12000 [00:14<00:00, 810.09it/s]
100%|| 12000/12000 [00:15<00:00, 781.09it/s]
100%|| 12000/12000 [00:15<00:00, 777.39it/s]
100%|| 12000/12000 [00:15<00:00, 754.40it/s]
100%|| 12000/12000 [00:20<00:00, 571.61it/s]

```

```

[689]: # Plot confusion matrix
M_mlp, acc_mlp = Confusion(X_test, y_test, mlpClassifier)
print ('MLP classifier accuracy: %f'%acc_mlp)
VisualizeConfussion(M_mlp)


```

```

100%|| 200/200 [00:00<00:00, 777.24it/s]

```

MLP classifier accuracy: 97.080000




```

[[0.98 0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.99 0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.97 0.  0.  0.  0.  0.  0.01 0. ]
 [0.  0.  0.01 0.97 0.  0.01 0.  0.  0.  0. ]

```

```
[0.  0.  0.  0.  0.99 0.  0.  0.  0.  0.01]
[0.  0.  0.  0.  0.  0.99 0.  0.  0.  0. ]
[0.01 0.  0.  0.  0.01 0.02 0.96 0.  0.  0. ]
[0.  0.  0.01 0.  0.01 0.  0.  0.95 0.  0.02]
[0.  0.  0.01 0.01 0.01 0.01 0.  0.  0.95 0. ]
[0.  0.  0.  0.01 0.03 0.  0.  0.  0.  0.96]]
```

```
[526]: # Plot filter weights
weights = mlpClassifier.weight1.data
plt.figure(figsize=(16,8))
for i in range(10):
    weight = weights[i,:]
    # normalize weights
    weight = normalize(weight)
    plt.subplot(2,5,i+1)
    plt.axis('off')
    plt.imshow(weight.reshape((28,28)), cmap="inferno")
```



Comments on weights: The weights do not look like as clearly as numbers like we observed in the previous case. Some of them look like number 2, some look like a combination of multiple numbers and some filters do not look like numbers at all.

Linear classifier was a special case where the output neuron is simply dot product with input image plus a bias and hence we had all filters looking like numbers. It may not be so here.

In above figure, that look somewhat like 2 (filter 5, filter 7, filter 8) points that the network is trying to fit different hidden neurons to the same hand-written digits with possibly different neuron activations for different strokes. But it's unclear what exactly other filters and if they represent anything tangible.

So, even though we see few patterns for the filters, we still have 100 hidden neurons and a ReLU non-linearity. It is very difficult to figure out what the neural network actually learns for each filter

at the hidden neurons because it has immense flexibility with the 100 units. This aspect essentially reflects in the 'hidden' part of the name 'hidden layer'. In conclusion, as networks grow deep and we keep adding non-linearities, the analysis of what network is doing becomes very difficult.

1.4.5 Problem 2.5: Convolutional Neural Network (CNN) [8 pts]

Here you will implement a CNN with the following architecture:

- `n=5`
- `ReLU(Conv(kernel_size=5x5, stride=2, output_features=n))`
- `ReLU(Conv(kernel_size=5x5, stride=2, output_features=n*2))`
- `ReLU(Linear(hidden units = 64))`
- `Linear(output_features=classes)`

So, 2 convolutional layers, followed by 1 fully connected hidden layer and then the output layer

Display the confusion matrix and accuracy after training. You should get around ~ 98 % accuracy for 10 epochs and batch size 50. **Note: You are not allowed to use `torch.nn.Conv2d()` and `torch.nn.Linear()`, Using these will lead to deduction of points. Use the declared `conv2d()`, `weight_variable()` and `bias_variable()` functions.** Although, in practice, when you move forward after this class you will use `torch.nn.Conv2d()` which makes life easier and hides all the operations underneath.

```
[696]: def conv2d(x, W, stride, bias=None):
    # x: input
    # W: weights (out, in, kH, kW)
    return F.conv2d(x, W, bias, stride=stride, padding=2)

# Defining a Convolutional Neural Network
class CNNClassifier(DNN):
    def __init__(self, classes=10, n=5):
        super(CNNClassifier, self).__init__()
        self.kernel_size = 5
        self.stride = 2
        self.ReLU = torch.nn.ReLU(inplace=True)
        self.n = n
        self.img_h = 28
        self.img_w = 28
        # Initialize weights and biases
        self.weight1 = weight_variable((self.n, 1, self.kernel_size, self.
→kernel_size))
        self.bias1 = bias_variable((self.n))
        self.weight2 = weight_variable(((self.n)**2, self.n, self.kernel_size, self.
→kernel_size))
        self.bias2 = bias_variable(((self.n)**2))
        self.weight3 = weight_variable((64, ((self.n)**2)*self.img_h//4*self.
→img_w//4))
        self.bias3 = bias_variable((64))
        self.weight4 = weight_variable((classes, 64))
```

```

        self.bias4 = bias_variable((classes))
        self.params = [self.weight1, self.bias1, self.weight2, self.bias2, self.
→weight3, self.bias3, self.weight4, self.bias4]

    def forward(self, x):
        x = torch.Tensor(x)
        # First convolution layer
        out1 = conv2d(x, self.weight1, stride=self.stride, bias=self.bias1)
        # Activation layer
        self.ReLU(out1)
        # Second convolution layer
        out2 = conv2d(out1, self.weight2, stride=self.stride, bias=self.bias2)
        # Activation layer
        self.ReLU(out2)
        # First linear layer
        out3 = torch.addmm(self.bias3, out2.view(list(out2.size())[0], -1), self.
→weight3.t())
        # Activation
        self.ReLU(out3)
        # Output layer
        y = torch.addmm(self.bias4, out3.view(list(out3.size())[0], -1), self.
→weight4.t())
        return y

cnnClassifier = CNNClassifier()
cnnClassifier.train_net(X_train, y_train, epochs=10)

```

```

100%|| 1200/1200 [00:08<00:00, 138.87it/s]
100%|| 1200/1200 [00:06<00:00, 175.41it/s]
100%|| 1200/1200 [00:06<00:00, 185.92it/s]
100%|| 1200/1200 [00:07<00:00, 167.84it/s]
100%|| 1200/1200 [00:06<00:00, 184.39it/s]
100%|| 1200/1200 [00:06<00:00, 184.99it/s]
100%|| 1200/1200 [00:06<00:00, 175.24it/s]
100%|| 1200/1200 [00:06<00:00, 185.21it/s]
100%|| 1200/1200 [00:06<00:00, 184.20it/s]
100%|| 1200/1200 [00:06<00:00, 185.00it/s]

```

```


[697]: # Plot confusion matrix
M_cnn, acc_cnn = Confusion(X_test, y_test, cnnClassifier)
print ('MLP classifier accuracy: %f'%acc_cnn)
VisualizeConfussion(M_cnn)

```

```

100%|| 200/200 [00:00<00:00, 541.80it/s]
MLP classifier accuracy: 98.280000

```



```
[0.99 0.  0.  0.  0.  0.  0.  0.  0.  0. ]
[0.  1.  0.  0.  0.  0.  0.  0.  0.  0. ]
[0.  0.  0.99 0.  0.  0.  0.  0.  0.  0. ]
[0.  0.  0.  0.99 0.  0.  0.  0.  0.  0. ]
[0.  0.  0.01 0.  0.98 0.  0.  0.  0.  0.01]
[0.  0.  0.  0.01 0.  0.99 0.  0.  0.  0. ]
[0.01 0.  0.  0.  0.  0.04 0.95 0.  0.  0. ]
[0.  0.  0.01 0.  0.  0.  0.  0.98 0.  0.01]
[0.01 0.  0.01 0.  0.  0.01 0.  0.  0.98 0. ]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.98]]
```

- Note that the MLP/ConvNet approaches lead to an accuracy a little higher than the K-NN approach.
- In general, Neural net approaches lead to significant increase in accuracy, but in this case since the problem is not too hard, the increase in accuracy is not very high.
- However, this is still quite significant considering the fact that the ConvNets we've used are relatively simple while the accuracy achieved using K-NN is with a search over 60,000 training images for every test image.

- You can look at the performance of various machine learning methods on this problem at <http://yann.lecun.com/exdb/mnist/>
- You can learn more about neural nets/ pytorch at https://pytorch.org/tutorials/beginner/deep_learning_6
- You can play with a demo of neural network created by Daniel Smilkov and Shan Carter at <https://playground.tensorflow.org/>